

Interactive voice-based direction giving agent

DIANA PERSICO and LUCIE GALLAND*

One of the drawbacks of voice-based direction giving is that the only solution for a user that did not understand the instruction or is not sure of the way to take is to look at the map. This option can however be dangerous and/or be a loss of time for the user. Since the real world can be different from maps, where some intersections are less ambiguous, and GPS positions are not always accurate, this situation is quite frequent. To bypass this issue, we propose to develop a direction-giving agent able to use grounding strategies to reduce this ambiguity. To this intent, we created a virtual world filled with ambiguous paths. Then, we collected data from six couples of human participants, pairing a first person who took the instructions to a second person who led the first one along a path. Based on this corpus we created a direction giving agent which at first gives a simple and direct instruction and then can reformulate it giving further pieces of information if asked. The performances of this interactive agent were compared with the ones of a simpler agent unable to give further information. We found that our interactive agent was better perceived and tend to have a better performance than a simple map-based navigation application.

Additional Key Words and Phrases: direction giving agent, grounding, interaction

ACM Reference Format:

Diana Persico and Lucie Galland. 2022. Interactive voice-based direction giving agent. 1, 1 (January 2022), 7 pages.

1 INTRODUCTION

Voice-based direction giving agents are widely used in common life by car drivers as well as by pedestrians or bike riders [Antikainen et al. 2006; Bartie and Mackaness 2006]. However, due to GPS tracking position errors and the ambiguity of real-life routes, simple direction giving agents can induce errors, wrong turns, and hesitations which can be dangerous and make people lose time [Boye et al. 2014]. A solution to avoid ambiguity is to look at the mini-map often provided with these agents. This solution can however be dangerous for all kinds of users: neither pedestrians, nor bike riders, nor car drivers are supposed to look at their phones [Looije et al. 2007]. On the other hand, those applications do not offer the possibility to interact with the agent and ask for clarifications such as “Do I turn there?” or “Could you repeat that?”. Not being able to interact with the user, these applications lack a central feature of human-human direction-giving: the conversational grounding process. As stated in [Clark and Schaefer 1989], the grounding criterion, along with content specification, is a necessary requirement for a contribution to the discourse. It’s by them defined as “The contributor and the

partners mutually believe that the partners have understood what the contributor meant to a criterion sufficient for current purposes.” The conversational grounding process makes the direction giving faster and reduces final errors. In this paper, we propose to add grounding strategies to a direction giving conversational agent, in order to reduce ambiguity and limit the number of wrong turns without losing time looking at a map. To this intent, we first collect and analyze a corpus of human-human interactions. Using the extracted strategies, we develop an agent able to reformulate with more detail its instruction if asked to do so. This agent is then tested on a virtual environment against a simple map-based application.

2 RELATED WORK

A rich part of the literature on direction giving agents argues on the usefulness of maps even in dangerous situations. Google navigation [goo [n. d.]], one of the most widely used navigation apps, uses spoken output and maps to guide pedestrians as if they were car drivers. Solutions were explored to improve the use of maps, Dobesova [Dobesova 2012] developed a voice-controlled map. However, a navigation application displaying maps on a small screen can be strenuous and confusing [Looije et al. 2007]. From this observation, [Boye et al. 2014] proposed a dialog-based navigation application. Their application uses grounding strategies to ensure that GPS positions, that can be noisy, are well interpreted. Another possibility for errors in such applications is incomprehension or misjudgment of the distance [Boye et al. 2014]. In [Edmonds 1993] authors present a computational model for a direction-giving model able to use grounding methods to ensure user’s comprehension. To this purpose, they start from a pre-planned direction giving dialog and adding grounding strategies along the conversation to achieve mutual belief that the plan has been understood. In [Baker et al. 2008] authors studied redundant utterances in case of a non-understanding signal during a direction giving task. To this intent, they asked a dyad of participants to reach a particular point in a building and then go back to the starting point. One participant gave directions while the other followed them. The interactions were recorded and transcribed. There were two conditions: one where the dyad could see each other and one where they could only hear each other. They found more redundant utterances in a direction giving task when the listener provided verbal or non-verbal signals of non-understanding in both conditions. Indeed, in direction giving tasks a final perfect understanding of the instructions is mandatory. This gives evidence of the importance of grounding in the direction-giving agent. They also found that most instructors’ instructions were reformulated in case of misunderstanding. A grounding direction giving agent should therefore be able to reformulate and improve its instructions in case of misunderstanding. A rich literature exists on how to improve direction giving application instructions. A large part of these works relies on landmarks. In [Michon and Denis 2001] authors have looked at the importance of landmarks in such dialogs. To this intent, they conducted two studies. The first one was composed of a participant that learned a route in the streets

*Both authors contributed equally to this research.

Authors’ address: Diana Persico; Lucie Galland.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2022 Association for Computing Machinery.

XXXX-XXXX/2022/1-ART \$15.00

<https://doi.org/>

of Paris. The participant then needed to guide a pedestrian through the same road, that was unknown to the pedestrian. They found that landmarks were used by guides more frequently at reorientation/critical nodes, specifically at points where several possible directions could be followed. This result showed the importance of landmarks in direction giving dialogs, even if only indirectly. To prove it more directly, they conducted a second study where participants were asked to follow simple instructions for the same route and then asked where these instructions were lacking. They found that the most common suggestion was the use of 3D landmarks at crucial points, proving their importance in the direction-giving dialogs. This study does not however address directly the fact that landmarks are not preferred to be used at all decision points. Indeed, in some cases, relative directions are more effective, as shown in [Götze and Boye 2015]. In this study, authors have asked to rate the confidence of pedestrians in relative or landmark-based instructions at different decision points. They found that in the case of simple decision points, participants were more confident in relative direction than in landmark-based direction. Based on this study, it is possible to determine which decision points should be indicated based on landmarks first hand and at which decision point landmarks should be introduced only in case of asked clarification. This study is done using a TTS and wizard of OZ setting, correcting another drawback of the first study, which only looked at human-human interaction. This indicates that even with voice assistants, the use of landmarks at complicated decision points is useful. One could argue that one reason that would make people prefer to use a map instead of interacting with a voice assistant is self-inhibition towards the use of voice assistants. As shown in [Easwara Moorthy and Vu 2014], users prefer to use voice assistants in a private location rather than in public places, such as in the streets when walking. This is a relevant factor to take into account, and as the authors found out, it could mostly depend on the participants' perceived social acceptability of the assistants. Nevertheless, inquiring on how to improve the performances of the assistants is an independent problem, and increasing their efficiency could lead users to overcome self-inhibition. We propose in this work to find appropriate grounding strategies for comprehension of the instruction of a direction giving agent and to test whether for this particular task a dialog system is more efficient than a map-based application. Hence, in this work, we answer the two following research questions:

- **RQ1:** What grounding strategies do humans use to improve direction giving efficiency?
- **RQ2:** Can grounding strategies replace the use of a map in a direction giving agent?

3 TECHNICAL DESCRIPTION AND METHODS

3.1 Path design

We designed a path with 25 turns (see fig.1). 20 of them are simple turns, where there is no ambiguity on the way to take. The remaining 5 are conceived to be tricky (an initial crossroads where you don't know which road you are facing; 2 different turns on a roundabout with 14 exits; 2 small alleys hard to notice). Landmarks were placed along the path in strategic points, to be used in the human-human interactions to give directions. We also add coins to collect along




Fig. 1. Path designed with landmarks

the path and time penalty for the use of the map to account for the role play bias demonstrated by [Ewald 2012]

3.2 Human-human interaction

To answer our **RQ1**, we collected a dataset of human-human interactions related to our particular path. This corpus is collected to identify points of interest in the path and which grounding strategies are used and efficient for humans in direction giving / following.

3.2.1 Collection setup. To create the text outputs of the assistant, we have first built a corpus of conversations of real people giving directions on our map. Six couples of participants (7 male and 5 female students of the École normale supérieure, all of them L2 English speakers) took part in the corpus data collection. We gave a copy of the map with the landmarks (fig.1) to each of the participants in the couple. We gave the map with the chosen path only to one of them, and we asked this person to guide the other one along this path. Since grounding is a process that can involve non-verbal behaviors, such as eye-gazes and head-nods [Nakano et al. 2003], and our goal was to collect conversations where grounding was only made verbally, we asked the participants to only look at their map and to not look at each other, so that they couldn't make use of non-verbal communication, which would have made verbal grounding poorer. The six conversations were audio-recorded and transcribed.

3.2.2 Interpretation of the corpus. To compare the conversations, we segmented the transcripts to make correspond each part of the text to a change of direction in the path. This comparison procedure allowed us first to confirm the points of interest, the tricky points which required more turns to be finally understood by the guided person. Then, to find out the most common grounding strategies used by 1. the direction-taker participants to show understanding, and 2. the direction-giver participants to clarify previous information after a request of the direction-taker participants for confirmation, reformulation, and further information. The landmarks were used a lot for the clarifications. When able to interact with another human giving direction, participants mainly gave backchannel feedback when the instruction was understood, and they asked


for reformulation or further information in case of ambiguity (see Table. 1). This allows us to answer our research question **RQ1**: The grounding strategies used by humans for efficient direction giving are mainly confirmation of understanding the instructions (about their present or next position on the map) and reformulation (more likely with addition of with further information). From the transcripts, it emerges that the direction-takers systematically used backchannels when the instruction was clear to them. The set of backchannels we found in the corpus is relatively small: okay, yeah, yes, uh-huh, right, sure, all right, yep, and a combination of those. In some cases, they used the repetition of the instruction as a backchannel. For our agent, we choose to use only the first strategy, using the set of backchannels we found, since the second one would have required language understanding. We then implemented an agent able to reformulate the direction-giving instructions with more details when asked.

3.3 Agent design

Our direction giving agent can give two versions of each instruction: a first simple instruction and a second more complex version that helps to reduce ambiguity. The original simple instruction is given by the agent every time the participant reaches a key point, while the more complex reformulation is provided when the user asks for it.

3.3.1 Instructions design. The original instructions and reformulations of the agents were computed from our corpus of human-human interactions. In here, we looked at the direction-givers' transcripts, and we selected over the first turn inputs of the participants (namely their first attempts of giving directions) some extracts with simple directions to be used for the first turn inputs of the agent. Then, we selected, cleaned (from repairs, repetitions, and other noises), and adjusted (a necessary step since our assistant is not able to understand semantic alignment or other more-refined strategies) the texts, to write the inputs of a further turn (the one coming after the participant's turn). Some examples of instructions and reformulations can be found in Table 2.

3.3.2 Implementation. The virtual environment was implemented using the Unity game engine. The generation and recognition of speech (Text To Speech and Speech To Text) are handled using Microsoft Azure software. Each time the user reaches a key point chosen in the previous section, the designed instruction is provided by the Azure TTS. A list of possible backchannels is computed using the human-human corpus collected previously. For every recognized speech, the agent checks if the input falls into the list of backchannels. The agent answers to a backchannel with another backchannel randomly picked from a list computed from the same dataset. Every recognized input that did not fall in the list of possible backchannels is considered as a request for reformulation or precision. In that case, the system provides the reformulated instruction designed in the previous section 3.3.1



(a) View of the participant in the game (b) View of the map by the participant

Fig. 2. Stimuli

4 EVALUATION

To answer our research question **RQ2**, a user study was conducted to measure the improvement in the performance and user's perception of the direction giving agent using our grounding strategy.

4.1 Stimuli

Participants were presented with a video game in which they were able to control a character (see fig. 2a). They were instructed on how to move in the game and how to consult the map, and they got brief access to a version of the simulation without the info about the path in a trial session. The trial session was a necessary step to familiarize with the simulated world for participants who were not used to playing video games. The participants were told that they needed to collect coins in a specific order and as fast as possible. To do so, they had to follow instructions from the direction giving agent. They were also able to consult a map that displayed the path they needed to take (see fig. 2b). Accessing the map gave a penalty of 4s to the participants where they were not able to move, being displayed with a black screen. We identified one between-subject independent variable (AgentType) which has two levels. First, the baseline (base type) only gives simple instructions and can not interact with the user. The second level (interactivetype) corresponds to the setting where participants can interact with the agent. They could ask for further information when needed. A video of an example interaction is available here.

4.2 Measurements

To measure the efficiency of the agent, we take into account the number of wrong turns, the time taken to complete the path, and the number of looks at the map. To measure their perception of the agent, participants were asked to rate their agreement from 1 to 5 with a list of statements (see Table.3), including participant satisfaction of the interaction (Satisfaction), whether the participant found the agent knowledgeable (Knowledgeable), whether the participant trusted the agent (Trust), whether the participant found useful to look at the map (Map), whether the participant found useful to listen to the agent's instructions (Listen), whether the participant would have liked to use the agent in the future (Future) and whether the participant would have recommended this agent to a friend (Recommend). To control with confounding variables, participants were also asked to provide their L1, their sex, their use and appreciation of assistants and their familiarity with videogames. The participants were all fluent in English, and we excluded the

Example	Talker	Utterance
Confirmation of understanding through backchannels		
1	DG	And at the first intersection, you go left.
	DT	Uh-huh. Okay.
2	DG	You take the middle one left
	DT	Middle one left, right.
Request for reformulation / further information		
3	DG	So you will have to take the sixth exit.
	DT	Ok wait so.
	DG	So basically-
	DT	I cross the planet? And after the planet, the second.
	DG	Yes so you cross the planet, you don't take one of the path that touches the planet.
	DT	Uh-huh.
	DG	you take the one just after it.

Table 1. Transcript extracts

Turn number	Original instruction	Reformulated instruction
1	You're the starting point. You go straight.	So, you take the path between the planet and the star.
12	You go left and you take the sixth exit.	So you go left, and there is a star, and after a cloud. And so you take the one after the cloud.
19	And after... it's the third left.	So, not the one going uphill: you take the one going downhill.
21	You continue and you take the second right.	So, it's the small alley that you need to take.

Table 2. Instructions examples

Measure	Question
Satisfaction	I am satisfied with my conversation experience
Knowledgeable	I found the agent knowledgeable
Trust	I trusted the agent
Map	I found useful to look at the map
Listen	I found useful to listen to the agent's instruction
Future	I would like to use this agent in the future
Recommend	I would recommend this agent to a friend

Table 3. Questionnaire on perception of the agent

language barrier as a confounding variable. As shown in [Weiss et al. 2003], the performances of women and men on specific tasks can be different: women tend to perform better than men on some verbal tasks, whereas men are better on some visual-spatial tasks. As in the article it is stated, it is hard to disentangle whether these differences in cognitive functions depend on biological sex or on gender issues, which could be more relevant [Napoleon 2008]. Participants were asked to answer questions about their use of assistants [Ashktorab et al. 2019] and their familiarity with videogames, rating from 1 to 5 their agreement with a list of statements (see Table.4), including whether participants frequently interact with virtual agents

(AssistantFreq), whether the participant enjoy talking to assistants (EnjoyAssistant), whether they find them useful (UsefulAssistant) and whether they are used to playing videogames (Videogames). To make the gap smaller between the participants used to videogames and the ones who were not used to, we made the trial session (see: Stimuli).

Measure	Question
AssistantFreq	How often do you interact with one or more virtual assistants (i.e. Google Assistant, Siri, Alexa, etc.)
EnjoyAssistant	I usually enjoy talking to assistant(s)
UsefulAssistant	The assistant(s) is/are useful
Videogames	I am used to play videogames

Table 4. Questionnaire on participants' habits

4.3 Hypotheses

We hypothesized the following:

- **H1:** The agent type (AgentType) of the conversational agent will have a main effect on the efficiency of the agent. More specifically, the participants when following the not interactive baseline agent's instructions (baselinetype) will be slower,

make more wrong turns, and look more at the map than the participants following the interactive agent's instructions (interactivetype).

- **H2:** The agent type (AgentType) of the conversational agent will have a main effect on the perceived quality of the interaction. More specifically, the agent using the not interactive baseline type (baselinetype) will be perceived as worse than the interactive agent (interactivetype).

4.4 Results

We collected data from 10 participants from different backgrounds. 40% of them were women, 60% were men. They were all fluent in English and originated from various countries. For 40% of them, their L1 was French, 50% originally spoke Italian, and 1 Farsi. Participants were assigned randomly to one of our 2 conditions.

4.4.1 Controls. To compute the assistant score, we first compute the Cronbachs alphas score on the 3 related questions (AssistantFreq, EnjoyAssistant, and UsefulAssistant) and found only acceptable reliability ($\alpha = 0,6$). However, using only the questions (AssistantFreq and EnjoyAssistant) we found good reliability ($\alpha = 0,77$). Therefore, for each participant, we compute the assistant score by averaging the score obtained for these 2 questions. A t-test showed that participants were significantly at the 95% level not equally distributed among the two conditions according to the assistant score ($p = 0.03$). In condition (interactivetype), the mean score is 2 and in condition (baselinetype) the mean is 3,7. Therefore, we need to control for this variable in our regression models.⁴

There is a difference on average in sex repartition over conditions, but the t-test showed that this difference was not significant ($p = 0.65$).

Finally, a t-test showed that the video game abilities differences across conditions are significant at the 90% level ($p = 0.06$). In condition (interactivetype), participants rated themselves on average at 4,5 while (baselinetype) participants rated themselves at 2,8.

Therefore, in our following regression models, it is necessary to account for the (videogames) and (assistant score) variables.

4.4.2 Efficiency. Our first metric for efficiency is the time taken to complete the path (see fig.3). A t-test shows that although the time taken is on average smaller for the interaction condition (488s against 570s), the difference is not significant ($p = 0,29$). To account for our control variables, we run a linear regression :

$$\text{time} = \alpha_1 \text{condition} + \alpha_2 \text{videogames} + \alpha_3 \text{assistant score} + \beta$$

where condition = 1 for (interactive type) and 0 for (baselinetype). The results are visible in Table 5. The p -value ($p = 0,4$) shows that the result is not significant. However, even with the control variables, it seems that participants in the condition (interactivetype) tend to take less time to complete the path than participants in the (baselinetype) condition. The (Condition) variable also has more impact than the (Videogames) or (Assistant score) variables.

A second metric we can use to measure efficiency is the number of wrong turns. (see fig.4). A t-test shows that although the number of wrong turns is on average smaller for the interactive condition (3,75s against 3,83s), the difference is however not significant ($p = 0,94$). To account for our control variables, we again run a linear regression




Fig. 3. Time taken to finish the path across conditions

Name of the variable	α	p -value
Condition	-108,9	0,4
Video game	16,5	0,64
Assistant score	0,08	0,90

Table 5. Linear regression on the total time

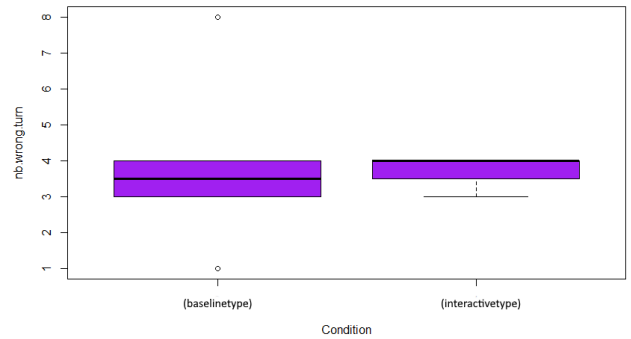


Fig. 4. Number of wrong turns across conditions

:

$$\text{wrongturns} = \alpha_1 \text{condition} + \alpha_2 \text{Videogames} + \alpha_3 \text{assistant score} + \beta$$

where condition = 1 for (interactive type) and 0 for (baselinetype). The results are visible in Table 6. The results are not significant, but the model predicts a small diminution of the number of wrong turns of approximately 0,16 with the (interactivetype) condition compared to the (baselinetype) condition. It seems that the video game abilities of the participant have more impact on the number of wrong turns. At the 90% confidence level, for each point of estimated video game ability, the participant make 0,8 fewer wrong turns.

A third and last measure for efficiency is the number of look at the map. (see fig.5). A t-test shows that the number of looks at

Name of the variable	α	p -value
Condition	-0,16	0,91
Video game	-0,79	0,1
Assistant score	-0,84	0,17

Table 6. Linear regression on the number of wrong turns




Fig. 5. Number of looks at the map across conditions

the map is significantly smaller at the 99% level for the interactive condition (0,25 against 5,83) with $p = 0,0007$. To account for our control variables, we also run a linear regression :

looks at the map = $\alpha_1 \text{condition} + \alpha_2 \text{Videogames} + \alpha_3 \text{assistant score} + \beta$

where condition = 1 for (interactive type) and 0 for (baselinetype). The results are visible in Table 7. With the control variables, the model predicts a significant diminution at the 99% level of looks at the map ($p = 0,01$). The diminution of the number of looks at the map is approximately 5,75 with the (interactivetype) condition compared to the (baselinetype) condition. The variable condition also has more impact than the two control variables. Due to our small

Name of the variable	α	p -value
Condition	-5,75	0,01
Video game	0,39	0,44
Assistant score	0,29	0,65

Table 7. Linear regression on the number of looks at the map

number of participants, we did not find significant effects for all our metrics. However, the results are coherent with our hypothesis **H1**: the time taken is on average smaller in the (interactive type) condition compared to the (baselinetype) condition. The number of wrong turns tends to be smaller in the (interactivetype) condition but seems more impacted by the control variable (Videogames). Finally, the number of looks at the map is significantly decreased by the (interactivetype) condition at the 99% confident level.




Fig. 6. Perception score across conditions

4.4.3 Perception score. To compute the perception score, we first compute the Cronbachs alphas score on the 7 questions about perception (Satisfaction, Knowledgeable, Map, Listen, Trust, Future, and Recommend) and found good reliability ($\alpha = 0,73$). For each participant, we compute the perception score by averaging the score obtained for each question, with the Map question counting as a negative score. A t-test showed a significant difference at the 95% level ($p = 0,05$) across conditions for the perception score (see fig.6) with a mean score of 0,9 for the (interactivetype) condition and of 0,7 for the (baselinetype) condition. To account for our control variables, we also run a linear regression :

perception score = $\alpha_1 \text{condition} + \alpha_2 \text{Videogames} + \alpha_3 \text{assistant score} + \beta$

where condition = 1 for (interactive type) and 0 for (baselinetype). The results are visible in Table 8. The linear regression shows a significant impact at the 99 % level of the condition on the perception score ($p = 0,007$). The (interactivetype) condition is perceived approximately 37% better than the (baseline type) condition, even when the control variables are taken into account. As expected, the assistant score also has a positive impact of approximately 10% significant at the 99% level ($p = 0,01$). Therefore our hypothesis **H2**

Name of the variable	α	p -value
Condition	0,37	0,007
Video game	0,009	0,71
Assistant score	0,11	0,01

Table 8. Linear regression on the perception of the agent

is verified.

5 CONCLUSION

In this article, we collected a small corpus of human-human interactions to determine the commonly used grounding strategies. From this corpus, we were able to answer our **RQ1**: reformulation is commonly used by humans to correct a misunderstanding of the guide. A direction-giving agent for a virtual environment was then created based on the before mentioned corpus. This agent can interact

with the user and reformulate its direction instruction when asked. We finally conducted a user study that answered our **RQ2**: Users perceive as better an interactive direction giving agent compared to a map-based navigation application. The results also tend towards showing that an interactive direction giving agent might be more efficient than a map-based application, but a more extensive study is necessary to significantly prove it.

5.1 Limitations

The game was not designed to handle large deviations from the path, which led to issues during the experiment when the participant did not follow the instructions. Some participants also could not finish the experiment because of motion sickness. Another limitation was that our agent is only able to reformulate once each instruction. The input of the participant is also classified in only two categories, the response of the agent could be more specific with a language understanding module. Another limitation was that many participants did not feel comfortable talking to the agent and preferred to look at the map, even though this was associated with a time penalty.

5.2 Future Works

In Future Work, we could cover the limitations by adding a language understanding module and more specific responses from the agent. The game could also be improved to avoid technical issues. Another way to improve our agent would be to refine the second grounding strategy emerging from the human-human interaction corpus: confirmation. A way to proceed could be to analyze the number of backchannels given by the user and react to a diminution of their frequency, which could signal misunderstanding or hesitation.

6 REPLICATION

Instruction and code necessary for replication can be found at <https://github.com/lgoue/Interactivedirection.git>

REFERENCES

- [n. d.]. Google Inc.: Google maps navigation. <http://www.google.com/mobile/navigation/>.
- Harri Antikainen, Jarmo Rusanen, Sami Vartiainen, Mauri Myllyaho, Jari Karvonen, Markku Oivo, Jouni Similä, and Kari Laine. 2006. Location-based Services as Tool for Developing Tourism in Marginal Regions. *Nordia Geographical Publications* 35 (12 2006), 39–50.
- Zahra Ashktorab, Mohit Jain, Q. Vera Liao, and Justin D. Weisz. 2019. *Resilient Chatbots: Repair Strategy Preferences for Conversational Breakdowns*. Association for Computing Machinery, New York, NY, USA, 1–12. <https://doi.org/10.1145/3290605.3300484>
- Rachel E. Baker, Alastair J. Gill, and Justine Cassell. 2008. Reactive Redundancy and Listener Comprehension in Direction-Giving. In *Proceedings of the 9th SIGdial Workshop on Discourse and Dialogue* (Columbus, Ohio) (*SIGdial '08*). Association for Computational Linguistics, USA, 37–45.
- Phil J. Bartie and William A. Mackaness. 2006. Development of a Speech-Based Augmented Reality System to Support Exploration of Cityscape. *Trans. GIS* 10 (2006), 63–86.
- Johan Boye, Morgan Fredriksson, Jana Götze, Joakim Gustafson, and Jürgen Königsmann. 2014. *Walk This Way: Spatial Grounding for City Exploration*. 59–67. https://doi.org/10.1007/978-1-4614-8280-2_6
- Herbert H. Clark and Edward F. Schaefer. 1989. Contributing to discourse. *Cognitive Science* 13, 2 (1989), 259–294. [https://doi.org/10.1016/0364-0213\(89\)90008-6](https://doi.org/10.1016/0364-0213(89)90008-6)
- Zdena Dobesova. 2012. Voice control of maps. 460–464. <https://doi.org/10.1109/TSP.2012.6256336>
- Aarthi Easwara Moorthy and Kim-Phuong Vu. 2014. Voice Activated Personal Assistant: Acceptability of Use in the Public Space. 324–334. https://doi.org/10.1007/978-3-319-07863-2_32
- Philip Edmonds. 1993. A Computational Model of Collaboration on Reference in Direction-Giving Dialogues.
- Jennifer D. Ewald. 2012. “can you tell me how to get there?”: Naturally-occurring versus role-play data in direction-giving. *Pragmatics* 22, 1 (2012), 79–102. <https://doi.org/10.1075/prag.22.1.03ewa>
- Jana Götze and Johan Boye. 2015. “Turn Left” Versus “Walk Towards the Café”: When Relative Directions Work Better Than Landmarks. https://doi.org/10.1007/978-3-319-16787-9_15
- Rosemarijn Looije, Guido Brake, and Mark Neerincx. 2007. Usability engineering for mobile maps. 532–539. <https://doi.org/10.1145/1378063.1378150>
- Pierre-Emmanuel Michon and Michel Denis. 2001. When and Why Are Visual Landmarks Used in Giving Directions?. In *Spatial Information Theory*, Daniel R. Montello (Ed.). Springer Berlin Heidelberg, Berlin, Heidelberg, 292–305.
- Yukiko Nakano, Gabe Reinstein, Tom Stocky, and Justine Cassell. 2003. Towards a Model of Face-to-Face Grounding. *Proceedings of the 41st Annual Meeting on Association for Computational Linguistics*, 553–561. <https://doi.org/10.3115/1075096.1075166>
- Siena Napoleon. 2008. From Here to There: A Sociolinguistic Study in Gender and Direction-Giving. *Indiana Undergraduate Journal of Cognitive Science* 2 (08 2008).
- Elisabeth M. Weiss, Georg Kemmler, Eberhard A. Deisenhammer, W.Wolfgang Fleischhacker, and Margarete Delazer. 2003. Sex differences in cognitive functions. (2003). [https://doi.org/10.1016/S0191-8869\(02\)00288-X](https://doi.org/10.1016/S0191-8869(02)00288-X)